

Natural Language Processing

Tutorial I: Regular Expressions and Text Normalization

Dr. Sun Aixin



Question Q1

- Write regular expressions for the following languages. By “word”, we mean an alphabetic string separated from other words by whitespace, any relevant punctuation, line breaks, and so forth. (**HINT**: please consult the book Chapter 2.1 or some websites on regular expressions.)
1. The set of all alphabetic strings;
 2. The set of all lower case alphabetic strings ending with a letter *b*;
 3. The set of all strings with two consecutive repeated words (e.g., “Humbert Humbert” and “the the” but not “the bug” or “the big bug”);
 4. All strings that start at the beginning of the line with an integer and that end at the end of the line with a word;
 5. All strings that have both the word *grotto* and the word *raven* in them (but not, e.g., words like *grottos* that merely *contain* the word *grotto*);

Question 1.1, 1.2

- The set of all alphabetic strings;
 - **$[a-zA-Z]^+$**
- The set of all lower case alphabetic strings ending with a letter b;
 - **$[a-z]^*b$**



Question 1.3

➤ The set of all strings with two consecutive repeated words (e.g., “Humbert Humbert” and “the the” but not “the bug” or “the big bug”);

- `(\b[a-zA-Z]+\b)\s+\1`

➤ Explanation

- `[a-zA-Z]+` → all alphabetic strings
- `\s` → whitespace (space, tab..)
- `\1` → used to refer to back to the first pattern in the expression which is put **inside a parentheses** `()`
 - We may have `\2` or `\3` to refer to the second and third patterns put inside parentheses.



Question 1.4

- All strings that start at the beginning of the line with an integer and that end at the end of the line with a word
 - `^\d+\b.*\b[a-zA-Z]+$`

➤ Explanation

- `\d` → a digit
- `\b` → a word boundary
- `^, $` → the **beginning** and **end** of a line
- `.` → a wildcard expression that matches any single character (except a carriage return)
- `*` → Kleene star, zero or more occurrences of the immediate previous character or regular expression
- `.*` → any string of characters



Question 1.5

- All strings that have both the word `grotto` and the word `raven` in them (but not, e.g., words like `grottos` that merely contain the word `grotto`)

- **`(.*\bgrotto\b.*\braven\b.*)"|(*\braven\b.*\bgrotto\b.*)"`**

- Explanation

- The two words `grotto` and `raven` may appear in any order.
 - There could be other strings around the two words

- <http://regexr.com/>

Question 2

- Try all your answers on <http://regexr.com/>
- You may need to change the textbox to test two cases: the textbox contains one or more matched strings, and the textbox does not contain any matched string.
- What are the errors (e.g., false positive and false negative) have you observed?



Question 2

- The set of all alphabetic strings;
 - **[a-zA-Z]+**

- The set of all lower case alphabetic strings ending with a letter b;
 - **[a-z]*b**

- The set of all strings with two consecutive repeated words (e.g., “Humbert Humbert” and “the the” but not “the bug” or “the big bug”);
 - **(\b[a-zA-Z]+\b)\s+\b**



Question 2

- All strings that start at the beginning of the line with an integer and that end at the end of the line with a word
 - `^\d+\b.*\b[a-zA-Z]+$`
- All strings that have both the word `grotto` and the word `raven` in them (but not, e.g., words like `grottos` that merely contain the word `grotto`)
 - `(.*\bgrotto\b.*\braven\b.*)|(.*\braven\b.*\bgrotto\b.*)`

Question 3

$$D[i, j] = \min \begin{cases} D[i-1, j] + 1 \\ D[i, j-1] + 1 \\ D[i-1, j-1] + \begin{cases} 1 & \text{if } source[i] \neq target[j] \\ 0 & \text{if } source[i] = target[j] \end{cases} \end{cases}$$

- Compute the edit distance (using **insertion cost 1**, **deletion cost 1**, **substitution cost 1**) of “idea” to “deal”. Show your work.

Target

	#	d	e	a	l
#	0	1	2	3	4
i	1				
d	2				
e	3				
a	4				

Source

	#	d	e	a	l
#	0	1	2	3	4
i	1	1			
d	2	1			
e	3	2			
a	4	3			

	#	d	e	a	l
#	0	1	2	3	4
i	1	1	2		
d	2	1	2		
e	3	2	1		
a	4	3	2		

Question 3

$$D[i, j] = \min \begin{cases} D[i-1, j] + 1 \\ D[i, j-1] + 1 \\ D[i-1, j-1] + \begin{cases} 1 & \text{if } source[i] \neq target[j] \\ 0 & \text{if } source[i] = target[j] \end{cases} \end{cases}$$

- Compute the edit distance (using **insertion cost 1**, **deletion cost 1**, **substitution cost 1**) of “idea” to “deal”. Show your work.

Target

	#	d	e	a	l
#	0	1	2	3	4
i	1	1	2		
d	2	1	2		
e	3	2	1		
a	4	3	2		

	#	d	e	a	l
#	0	1	2	3	4
i	1	1	2	3	
d	2	1	2	3	
e	3	2	1	2	
a	4	3	2	1	


	#	d	e	a	l
#	0	1	2	3	4
i	1	1	2	3	4
d	2	1	2	3	4
e	3	2	1	2	3
a	4	3	2	1	2

Source

Question 3

$$D[i, j] = \min \begin{cases} D[i-1, j] + 1 \\ D[i, j-1] + 1 \\ D[i-1, j-1] + \begin{cases} 1 & \text{if } source[i] \neq target[j] \\ 0 & \text{if } source[i] = target[j] \end{cases} \end{cases}$$

- Compute the edit distance (using **insertion cost 1**, **deletion cost 1**, **substitution cost 1**) of “idea” to “deal”. Show your work.



Question 4

$$D[i, j] = \min \begin{cases} D[i-1, j] + 1 \\ D[i, j-1] + 1 \\ D[i-1, j-1] + \begin{cases} 1 & \text{if } source[i] \neq target[j] \\ 0 & \text{if } source[i] = target[j] \end{cases} \end{cases}$$

- Compute the edit distance (using **insertion cost 1**, **deletion cost 1**, **substitution cost 1**) of two sentences “compute the edit distance” to “the edit distance is computed”. Show your work and show the alignment between the two strings.
- If similar questions appear in exam, the requirement will be made clear, whether the edit distance is to be computed at character level or at word level.
 - We use the first character to represent each word, for simplicity and clarity.

		Target					
		#	T	E	D	I	C
Source	#	0	1	2	3	4	5
	C	1					
	T	2					
	E	3					
	D	4					

	#	T	E	D	I	C
#	0	1	2	3	4	5
C	1	1	2	3	4	5
T	2	1	2	3	4	5
E	3	2	1	2	3	4
D	4	3	2	1	2	3

Question 4

$$D[i, j] = \min \begin{cases} D[i-1, j] + 1 \\ D[i, j-1] + 1 \\ D[i-1, j-1] + \begin{cases} 1 & \text{if } source[i] \neq target[j] \\ 0 & \text{if } source[i] = target[j] \end{cases} \end{cases}$$

- Compute the edit distance (using **insertion cost 1**, **deletion cost 1**, **substitution cost 1**) of two sentences “compute the edit distance” to “the edit distance is computed”. Show your work and show the alignment between the two strings.

compute **the edit distance**

the edit distance is computed

Target

Source		#	T	E	D	I	C
	#	0	1	2	3	4	5
	C	1	1	2	3	4	5
	T	2	1	2	3	4	5
	E	3	2	1	2	3	4
	D	4	3	2	1	2	3

	#	T	E	D	I	C
#	0	1	2	3	4	5
C	1	1	2	3	4	5
T	2	1	2	3	4	5
E	3	2	1	2	3	4
D	4	3	2	1	2	3

Natural Language Processing

Regular Expressions and FSA

**FSA is non-examinable.
For your information only.**

Dr. Sun Aixin



Regular Expression and Finite state automata (FSA)

➤ Regular expressions

- Regex is compact textual strings (e.g., “/[tT]he/”), which is perfect for specifying patterns in programs or command-lines
- Regex can be implemented as an FSA (Finite State Automata)

➤ Finite state automata

- Graphs: **Nodes** are states and **edges** are transitions among states
- FSA has a wide range of uses

➤ FSA vs Regex


- FSA can be described with a regular expression
- Regular expression is a textual way of specifying the structure of FSA



FSA as Graphs

FSA is non-examinable


- Example regular expression: `/baa+!/`
- Corresponding FSA



It has 5 states (q_0 to q_4)
b, **a**, and **!** are in its alphabet
 q_0 is the start state
 q_4 is an accept state
It has 5 transitions


- A set of states: Q
 - A finite alphabet: Σ
 - A start state
 - A set of accept/final states
 - A transition function that maps $Q \times \Sigma \rightarrow Q$
- FSA**

Don't take term **alphabet** word too narrowly;
it just means we need a **finite set of symbols** in the input.



Yet Another View


- An FSA can ultimately be represented as a **table**



If currently in state q_1 and an input is **a**, then go to state q_2

	Input		
State	b	a	!
q_0	1		
q_1		2	
q_2		3	
q_3		3	4
q_4 :			




- **Recognition** is the process of determining if a string should be accepted by a machine



Recognition (D-Recognize)

➤ D-Recognize (in a tap-view)


- Starting the process from the **start state**
- Examining the **current input**
- Consulting **the table**
- Going **to a new state** and updating the tape pointer.
- Until you run **out of tape**. Accept?



	Input		
State	b	a	!
q_0	1		
q_1		2	
q_2		3	
q_3		3	4
q_4 :			


Deterministic FSA vs Non-Deterministic FSA

- If currently at a state, given an input
- Deterministic FSA: There is only one next state to move to
 - Non-deterministic FSA: There are more than one possible states to move to




Deterministic FSA vs Non-Deterministic FSA

- There exists other form of Non-deterministic FSA
 - Epsilon transitions
 - These transitions do not examine or advance the tape during recognition




Deterministic FSA vs Non-Deterministic FSA

- Non-deterministic machines can be **converted** to deterministic through algorithm
- Non-deterministic machines **are not more powerful** than deterministic ones in terms of the languages they can and can't characterize
 - Not always obvious to users whether the regex that they've produced is deterministic or non-deterministic
 - Sometimes, non-determinism may look more natural (understandable)




Additional Question on FSA

- Write a regular expression for the language accepted by the following NFSA.




Hint: List the strings can be generated from the NFSA



- ab
- aba
- abab
- ababa
- ababab
- abaab
- ...

Answer

➤ $(aba^+)^+$



$q_0 ::= \varepsilon$

$q_1 ::= q_0a \mid q_2a$

$q_1 ::= a \mid q_2a$

$q_2 ::= q_1b \mid q_3a$

$q_3 ::= q_1b$

$q_2 ::= q_1b \mid q_1ba$

$q_2 ::= ab \mid q_2ab \mid aba \mid q_2aba$

$q_2 ::= ab \mid aba \mid q_2ab \mid q_2aba$

$q_2 ::= (aba^+)|(q_2aba^+)$

$q_2 ::= (aba^+)^+$